

Lab 2

Matrix Multiplication (Multi-Threading)



Name : Pancee Wahid

ID: 18010467

Code Overview

```
C main.c ×
C main.c > ...
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <pthread.h>
4  #include <string.h>
5  #include <sys/time.h>
6
7  typedef struct {
8      int** data;
9      int r;
10     int c;
11 } Matrix;
12
13 typedef struct {
14     Matrix* mat_1;
15     Matrix* mat_2;
16     Matrix* result;
17     int row;
18     int col;
19 } ThreadData;
20
21 Matrix* matrix_mem_alloc(int r, int c);
22 Matrix* read_matrix(char* input);
23 void write_file(Matrix* result, char* file_name, int method);
24 void thread_per_matrix(Matrix* result, Matrix* mat_1, Matrix* mat_2);
25 void* thread_per_row(void* arg);
26 void* thread_per_element(void* arg);
27
28
29 > int main(int argc, char* argv[]){ ...
144
145 > Matrix* matrix_mem_alloc(int r, int c){ ...
154
155 > Matrix* read_matrix(char* input){ ...
184
185 > void write_file(Matrix* result, char* file_name, int method){ ...
223
224 > void thread_per_matrix(Matrix* result, Matrix* mat_1, Matrix* mat_2){ ...
233
234 > void* thread_per_row(void* arg){ ...
245
246 > void* thread_per_element(void* arg){ ...
```

Structs:

→ Matrix

- data: a double pointer to the matrix
- r: number of columns of the matrix
- c: number of rows of the matrix

→ ThreadData

- mat_1: the first matrix double pointer
- mat_2: the second matrix double pointer
- result: double pointer to the matrix to store the result in
- row: row to be used (used in thread per row and thread per element methods)
- col: col to be used (used in thread per element method)

Functions

Main Function

```
29  int main(int argc, char* argv[]){
30      char* input_1 = (char*)malloc(100*sizeof(char));
31      char* input_2 = (char*)malloc(100*sizeof(char));
32      char* output = (char*)malloc(100*sizeof(char));
33      struct timeval start, stop;
34
35      // Handling files names
36      if (argc == 1){ // no file names are entered
37          input_1 = "a.txt";
38          input_2 = "b.txt";
39          output = "c";
40      }
41      else if (argc == 4){ // file names are entered
42          strcpy(input_1, argv[1]);
43          strcat(input_1, ".txt");
44
45          strcpy(input_2, argv[2]);
46          strcat(input_2, ".txt");
47
48          strcpy(output, argv[3]);
49      }
50      else {
51          printf("Invalid Arguments!");
52          exit(1);
53      }
54
55      // Reading matrix A and B from the specified files
56      Matrix* mat_1 = read_matrix(input_1);
57      Matrix* mat_2 = read_matrix(input_2);
58      free(input_1);
59      free(input_2);
60
61      // Check compatibility
62      if (mat_1->c != mat_2->r){
63          printf("Incompatible sizes!\n");
64          exit(1);
65      }
66
67      // Apply method 1 : Thread per matrix
68      Matrix* result_1 = matrix_mem_alloc(mat_1->r, mat_2->c);
69      gettimeofday(&start, NULL);
70      thread_per_matrix(result_1, mat_1, mat_2);
71      gettimeofday(&stop, NULL);
72      char* out_file_1 = (char*)malloc(100*sizeof(char));
73      strcpy(out_file_1, output);
74      strcat(out_file_1, "_per_matrix.txt");
75      write_file(result_1, out_file_1, 1);
76      printf("Microseconds taken by thread by matrix: %lu\n", stop.tv_usec - start.tv_usec);
77      free(result_1);
78      free(out_file_1);
```

```

80 // Apply method 2 : Thread per row
81 Matrix* result_2 = matrix_mem_alloc(mat_1->r, mat_2->c);
82 gettimeofday(&start, NULL);
83 pthread_t rows_thread[mat_1->r]; // declare threads
84 for (int i = 0; i < mat_1->r; i++){
85     // pack data needed in struct
86     ThreadData* thread_row_data = (ThreadData *)malloc(sizeof(ThreadData));
87     thread_row_data->mat_1 = mat_1;
88     thread_row_data->mat_2 = mat_2;
89     thread_row_data->result = result_2;
90     thread_row_data->row = i;
91     // create thread for row i
92     pthread_create(&rows_thread[i], NULL, thread_per_row, (void*)(thread_row_data));
93 }
94 // join (wait) the rows threads
95 for (int i = 0; i < mat_1->r; i++) {
96     pthread_join(rows_thread[i], NULL);
97 }
98 gettimeofday(&stop, NULL);
99 char* out_file_2 = (char*)malloc(100*sizeof(char));
100 strcpy(out_file_2, output);
101 strcat(out_file_2, "_per_row.txt");
102 write_file(result_2, out_file_2, 2);
103 printf("Microseconds taken by thread by row: %lu\n", stop.tv_usec - start.tv_usec);
104 free(result_2);
105 free(out_file_2);

```

```

107 // Apply method 3 : Thread per element
108 Matrix* result_3 = matrix_mem_alloc(mat_1->r, mat_2->c);
109 gettimeofday(&start, NULL);
110 pthread_t elements_thread[mat_1->r][mat_2->c]; // declare threads
111 for (int i = 0; i < mat_1->r; i++){
112     for (int j = 0; j < mat_2->c; j++){
113         // pack data needed in struct
114         ThreadData* thread_element_data = (ThreadData *)malloc(sizeof(ThreadData));
115         thread_element_data->mat_1 = mat_1;
116         thread_element_data->mat_2 = mat_2;
117         thread_element_data->result = result_2;
118         thread_element_data->row = i;
119         thread_element_data->col = j;
120         // create thread for element[i][j]
121         pthread_create(&elements_thread[i][j], NULL,
122             thread_per_element, (void*)(thread_element_data));
123     }
124 }
125 // join (wait) the elements threads
126 for (int i = 0; i < mat_1->r; i++) {
127     for (int j = 0; j < mat_2->c; j++)
128         pthread_join(elements_thread[i][j], NULL);
129 }
130 gettimeofday(&stop, NULL);
131 char* out_file_3 = (char*)malloc(100*sizeof(char));
132 strcpy(out_file_3, output);
133 strcat(out_file_3, "_per_element.txt");
134 write_file(result_3, out_file_3, 3);
135 printf("Microseconds taken by thread by element: %lu\n", stop.tv_usec - start.tv_usec);
136 free(result_3);
137 free(out_file_3);
138
139 free(mat_1);
140 free(mat_2);
141 free(output);
142
143 return 0;
144 }

```

matrix_mem_alloc():

It takes the number of rows and columns of the matrix and allocates memory for them then returns a pointer to the struct Matrix created for the matrix holding a double pointer to it, its number of rows and its number of columns.

```
146 Matrix* matrix_mem_alloc(int r, int c){
147     Matrix* mat = (Matrix*)malloc(sizeof(Matrix*));
148     mat->r = r;
149     mat->c = c;
150     mat->data = (int**)malloc(r * sizeof(int*));
151     for (int i = 0; i < r; i++)
152         mat->data[i] = (int*)malloc(c * sizeof(int));
153     return mat;
154 }
```

read_matrix():

It takes the name of the file that contains the matrix to be read, reads its number of rows and columns and calls matrix_mem_alloc() to allocate memory for the read matrix. Then it reads the matrix elements storing them in their locations and returns a pointer to the Matrix created.

```
156 Matrix* read_matrix(char* input){
157     FILE* file;
158     int r,c;
159     file = fopen(input, "r");
160     if (file == NULL){
161         printf("Can't open file %s\n", input);
162         exit(1);
163     }
164
165     if(fscanf(file, "row=%d col=%d", &r, &c) != 2){
166         printf("Invalid format in file %s. Can't extract number of rows and columns!\n", input);
167         exit(1);
168     }
169
170     Matrix* matrix = matrix_mem_alloc(r, c);
171     for (int i = 0; i < matrix->r; i++){
172         for (int j = 0; j < matrix->c; j++){
173             if(fscanf(file, "%d" ,&matrix->data[i][j]) != 1){
174                 printf("Invalid format in file %s. Can't read the matrix!\n", input);
175                 free(matrix);
176                 exit(1);
177             }
178         }
179     }
180
181     fclose(file);
182
183     return matrix;
184 }
```

write_file():

It takes the required data to store the resulting matrix from multiplication of the two read matrices in the specified file.

```
186 void write_file(Matrix* result, char* file_name, int method){
187     FILE* file;
188     file = fopen(file_name, "a");
189
190     // print error msg in case of error while creating the file
191     if (file == NULL){
192         printf("Error in creating file!");
193         exit(1);
194     }
195
196     switch (method){
197         case 1:
198             fprintf(file, "Method: A thread per matrix\n");
199             break;
200         case 2:
201             fprintf(file, "Method: A thread per row\n");
202             break;
203         case 3:
204             fprintf(file, "Method: A thread per element\n");
205             break;
206         default:
207             printf("Error in printing! Invalid method number!\n");
208             exit(1);
209     }
210     fprintf(file, "row=%d col=%d\n", result->r, result->c);
211     int i, j;
212     for (i = 0; i < result->r; i++){
213         for (j = 0; j < result->c - 1; j++){
214             fprintf(file, "%d ", result->data[i][j]);
215
216             if (i < result->r - 1)
217                 fprintf(file, "%d\n", result->data[i][j]);
218             else
219                 fprintf(file, "%d", result->data[i][j]);
220         }
221
222         fclose(file); // close the file
223     }
```

thread_per_matrix():

It calculates the product **matrix** of multiplying matrix 1 by matrix 2 normally by multiplying every row in matrix 1 by every column in matrix2.

thread_per_row():

It calculates the specified **row** of the product matrix by multiplying the row with the same index from matrix 1 with every column of matrix 2 storing the results in the result matrix.

thread_per_element():

It calculates the specified **element** of the product matrix by multiplying the row with the same index from matrix 1 with the column of the same index of matrix 2 storing the results in the result matrix.

```
225 void thread_per_matrix(Matrix* result, Matrix* mat_1, Matrix* mat_2){
226     for (int i = 0; i < result->r; i++){ //for each row
227         for (int j = 0; j < result->c; j++){ // for each col
228             result->data[i][j] = 0;
229             for (int k = 0; k < mat_1->c; k++)
230                 result->data[i][j] += (mat_1->data[i][k] * mat_2->data[k][j]);
231         }
232     }
233 }
234
235 void* thread_per_row(void* arg){
236     ThreadData* data = (ThreadData*) arg;
237     // multiply row i of mat_1 by each column of mat_2 forming row i of result
238     for (int j = 0; j < data->result->c; j++){
239         data->result->data[data->row][j] = 0;
240         for (int k = 0; k < data->mat_1->c; k++){
241             data->result->data[data->row][j] +=
242                 (data->mat_1->data[data->row][k] * data->mat_2->data[k][j]);
243         }
244     }
245     free(data);
246     pthread_exit(NULL);
247 }
248
249 void* thread_per_element(void* arg){
250     ThreadData* data = (ThreadData*) arg;
251     data->result->data[data->row][data->col] = 0;
252     for (int k = 0; k < data->mat_1->c; k++){
253         data->result->data[data->row][data->col] +=
254             (data->mat_1->data[data->row][k] * data->mat_2->data[k][data->col]);
255     }
256     free(data);
257     pthread_exit(NULL);
258 }
259
260 }
```

Sample Runs

The given 3 test cases give the following output:

```
pancee79@ubuntu: ~/Desktop/Operating_Systems/Matrix_Multiplication
File Edit View Search Terminal Help
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ gcc main.c -o matMultp -lpthread
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ ./matMultp 1_a 1_b 1
Microseconds taken by thread by matrix: 5
Microseconds taken by thread by row: 1168
Microseconds taken by thread by element: 10126
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ ./matMultp 2_a 2_b 2
Microseconds taken by thread by matrix: 0
Microseconds taken by thread by row: 611
Microseconds taken by thread by element: 1120
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ ./matMultp 3_a 3_b 3
Incompatible sizes!
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$
```

Test case 3

Input matrices

3_a.txt					3_b.txt			
row=5 col=5					row=4 col=4			
1	2	3	4	5	1	2	3	4
6	7	8	9	10	5	6	7	8
11	12	13	14	15	9	10	11	12
16	17	18	19	20	13	14	15	16
21	22	23	24	25				

Since these matrices can't be multiplied → as they don't follow the rule that if size of A is $n \times m$, then size of B should be $m \times k$ (number of columns of first matrix \neq number of rows of second matrix)
Therefore, the program printed an error message that size is incompatible and terminated.

Test case 1

Input matrices

Open ▾ 1_a.txt ~./Des... Save ≡ ⌵ ⌵ ⌵

row=10 col=5
1 2 3 4 5
6 7 8 9 10
11 12 13 14 15
16 17 18 19 20
21 22 23 24 25
26 27 28 29 30
31 32 33 34 35
36 37 38 39 40
41 42 43 44 45
46 47 48 49 50

Tab Width: 8 ▾ Ln 11, Col 35 ▾ INS

Open ▾ 1_b.txt ~./Desktop/Operating_Systems/Matrix_Multiplication Save ≡ ⌵ ⌵ ⌵

row=5 col=10
1 2 3 4 5 6 7 8 9 10
11 12 13 14 15 16 17 18 19 20
21 22 23 24 25 26 27 28 29 30
31 32 33 34 35 36 37 38 39 40
41 42 43 44 45 46 47 48 49 50

Plain Text ▾ Tab Width: 8 ▾ Ln 6, Col 75 ▾ INS

Output matrix

Open ▾ 1_per_matrix.txt ~./Desktop/Operating_Syst... Save ≡ ⌵ ⌵ ⌵

1_per_matrix.txt × 1_per_row.txt × 1_per_element.txt ×
Method: A thread per matrix
row=10 col=10
415 430 445 460 475 490 505 520 535 550
940 980 1020 1060 1100 1140 1180 1220 1260 1300
1465 1530 1595 1660 1725 1790 1855 1920 1985 2050
1990 2080 2170 2260 2350 2440 2530 2620 2710 2800
2515 2630 2745 2860 2975 3090 3205 3320 3435 3550
3040 3180 3320 3460 3600 3740 3880 4020 4160 4300
3565 3730 3895 4060 4225 4390 4555 4720 4885 5050
4090 4280 4470 4660 4850 5040 5230 5420 5610 5800
4615 4830 5045 5260 5475 5690 5905 6120 6335 6550
5140 5380 5620 5860 6100 6340 6580 6820 7060 7300

Plain Text ▾ Tab Width: 8 ▾ Ln 12, Col 50 ▾ INS

Open ▾ 1_per_row.txt ~./Desktop/Operating_Syst... Save ≡ ⌵ ⌵ ⌵

1_per_matrix.txt × 1_per_row.txt × 1_per_element.txt ×
Method: A thread per row
row=10 col=10
415 430 445 460 475 490 505 520 535 550
940 980 1020 1060 1100 1140 1180 1220 1260 1300
1465 1530 1595 1660 1725 1790 1855 1920 1985 2050
1990 2080 2170 2260 2350 2440 2530 2620 2710 2800
2515 2630 2745 2860 2975 3090 3205 3320 3435 3550
3040 3180 3320 3460 3600 3740 3880 4020 4160 4300
3565 3730 3895 4060 4225 4390 4555 4720 4885 5050
4090 4280 4470 4660 4850 5040 5230 5420 5610 5800
4615 4830 5045 5260 5475 5690 5905 6120 6335 6550
5140 5380 5620 5860 6100 6340 6580 6820 7060 7300

Plain Text ▾ Tab Width: 8 ▾ Ln 1, Col 1 ▾ INS

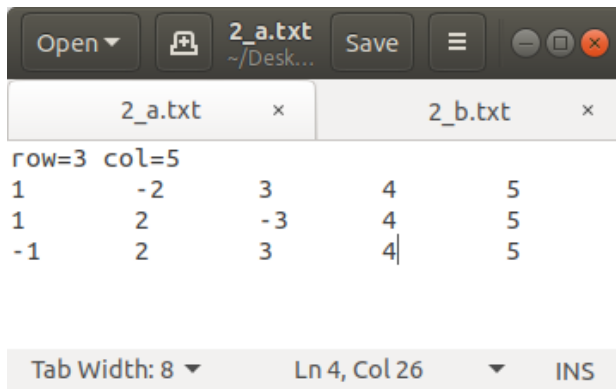
Open ▾ 1_per_element.txt ~./Desktop/Operating_Syst... Save ≡ ⌵ ⌵ ⌵

1_per_matrix.txt × 1_per_row.txt × 1_per_element.txt ×
Method: A thread per element
row=10 col=10
415 430 445 460 475 490 505 520 535 550
940 980 1020 1060 1100 1140 1180 1220 1260 1300
1465 1530 1595 1660 1725 1790 1855 1920 1985 2050
1990 2080 2170 2260 2350 2440 2530 2620 2710 2800
2515 2630 2745 2860 2975 3090 3205 3320 3435 3550
3040 3180 3320 3460 3600 3740 3880 4020 4160 4300
3565 3730 3895 4060 4225 4390 4555 4720 4885 5050
4090 4280 4470 4660 4850 5040 5230 5420 5610 5800
4615 4830 5045 5260 5475 5690 5905 6120 6335 6550
5140 5380 5620 5860 6100 6340 6580 6820 7060 7300

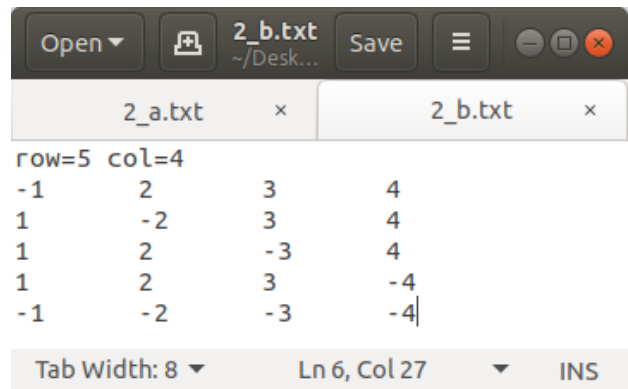
Plain Text ▾ Tab Width: 8 ▾ Ln 3, Col 32 ▾ INS

Test case 2

Input matrices

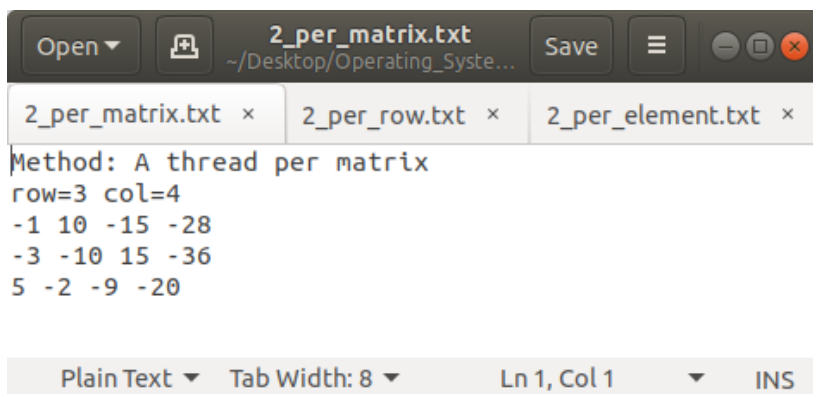


```
2_a.txt x 2_b.txt x
row=3 col=5
1 -2 3 4 5
1 2 -3 4 5
-1 2 3 4 5
Tab Width: 8 Ln 4, Col 26 INS
```

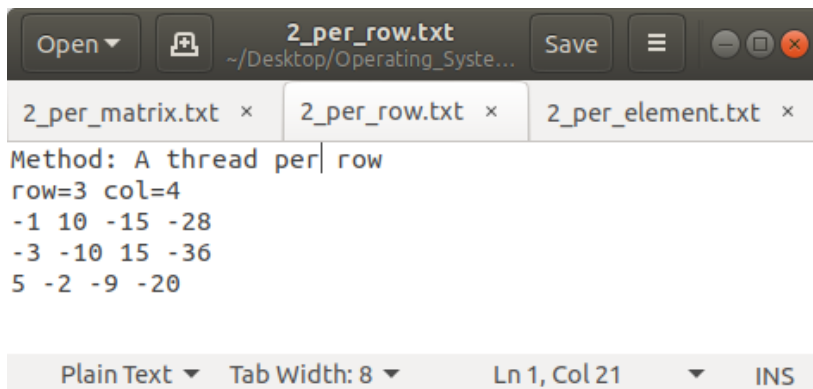


```
2_a.txt x 2_b.txt x
row=5 col=4
-1 2 3 4
1 -2 3 4
1 2 -3 4
1 2 3 -4
-1 -2 -3 -4
Tab Width: 8 Ln 6, Col 27 INS
```

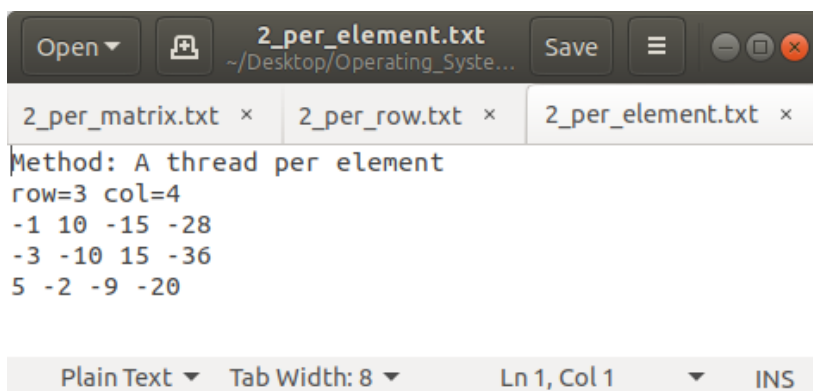
Output matrix



```
2_per_matrix.txt x 2_per_row.txt x 2_per_element.txt x
Method: A thread per matrix
row=3 col=4
-1 10 -15 -28
-3 -10 15 -36
5 -2 -9 -20
Plain Text Tab Width: 8 Ln 1, Col 1 INS
```



```
2_per_row.txt x 2_per_element.txt x
Method: A thread per row
row=3 col=4
-1 10 -15 -28
-3 -10 15 -36
5 -2 -9 -20
Plain Text Tab Width: 8 Ln 1, Col 21 INS
```



```
2_per_element.txt x
Method: A thread per element
row=3 col=4
-1 10 -15 -28
-3 -10 15 -36
5 -2 -9 -20
Plain Text Tab Width: 8 Ln 1, Col 1 INS
```

Comparison

```
pancee79@ubuntu: ~/Desktop/Operating_Systems/Matrix_Multiplication
File Edit View Search Terminal Help
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ gcc main.c -o matMultp -lpthread
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ ./matMultp 1_a 1_b 1
Microseconds taken by thread by matrix: 5
Microseconds taken by thread by row: 1168
Microseconds taken by thread by element: 10126
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ ./matMultp 2_a 2_b 2
Microseconds taken by thread by matrix: 0
Microseconds taken by thread by row: 611
Microseconds taken by thread by element: 1120
```

```
pancee79@ubuntu: ~/Desktop/Operating_Systems/Matrix_Multiplication
File Edit View Search Terminal Help
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ ./matMultp med_a med_b med
Microseconds taken by thread by matrix: 208
Microseconds taken by thread by row: 1400
Microseconds taken by thread by element: 49323
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ ./matMultp big_a big_b big
Microseconds taken by thread by matrix: 1692
Microseconds taken by thread by row: 3751
Microseconds taken by thread by element: 233164
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$ ./matMultp bigger_a bigger_b bigger
Microseconds taken by thread by matrix: 22870
Microseconds taken by thread by row: 18478
Microseconds taken by thread by element: 18446744073708838248
pancee79@ubuntu:~/Desktop/Operating_Systems/Matrix_Multiplication$
```

Note:

- The used matrices for med_a and med_b were of size (20x30) and (30x50).
- The used matrices for big_a and big_b were of size (40x60) and (60x100).
- The used matrices for bigger_a and bigger_b were of size (150x150) and (150x150).

Can be found [here](#).

Observation:

- For the test cases shown before, med and big matrices, performing matrix multiplication with one thread for the whole matrix was faster than the two other methods.
- For bigger matrices, performing matrix multiplication with one thread per row was better than the two other methods.
- The method of creating thread to calculate each element was bad in small matrices and worst in bigger ones.

Conclusion:

After running the three methods: thread per matrix, thread per row and thread per element on matrices of different sizes, it was deduced that:

- For small matrices (generally not too many computations), it's more efficient to use a single thread to avoid the overhead of creating threads which will waste much time than performing the computations using a single thread.
- For large matrices (generally heavy computations), it's more efficient to distribute the work on a reasonable number of threads which will improve the execution time of the program. However, creating so many threads may be more waste of time and resources than not using threads at all.